



# §3.3 MLBD MRes practical

## From a single neuron to the multilayer perceptron

Jarvist Moore Frost.



Twitter: @JarvistFrost

Imperial College London  
Email: [jarvist.frost@ic.ac.uk](mailto:jarvist.frost@ic.ac.uk)  
<https://jarvist.github.io>

# Intended Learning Outcomes §3.3

Demonstrate how multilayer perceptrons are built, by building a network and solving a task within PyTorch.

- Be able to summarise what we learn in the first session. (Taught from these slides)
- Reimplement a single neuron binary classifier in PyTorch. (Jupyter)
- Be aware of the mathematics of backpropagation, the multi-layer learning algorithm. (Jupyter)
- Use a multilayer perceptron to solve a non-linear task. (Jupyter)
- Run well defined machine-learning experiments to explore and document the parameter space of model construction and learning, communicating progress with your peers. (Jupyter)

# Recommended reading

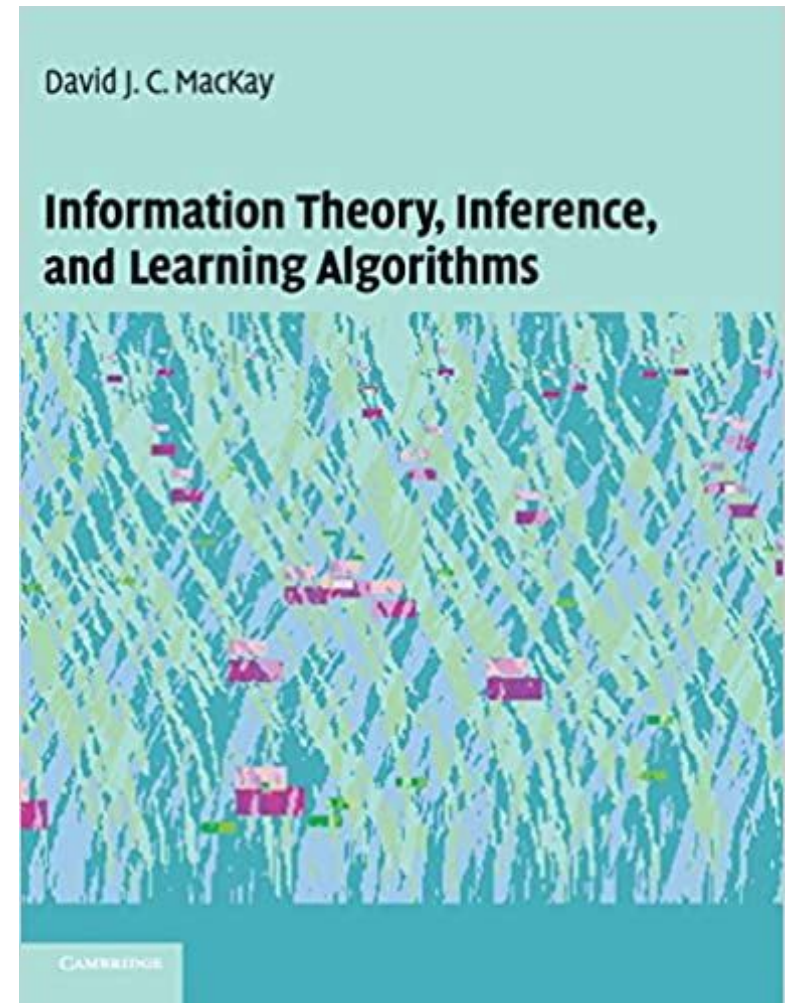
David MacKay, Information Theory, Inference and Learning Algorithms (ITILA), 2003, Chapters 38–42.

Freely available online!

A physicists perspective on ML.

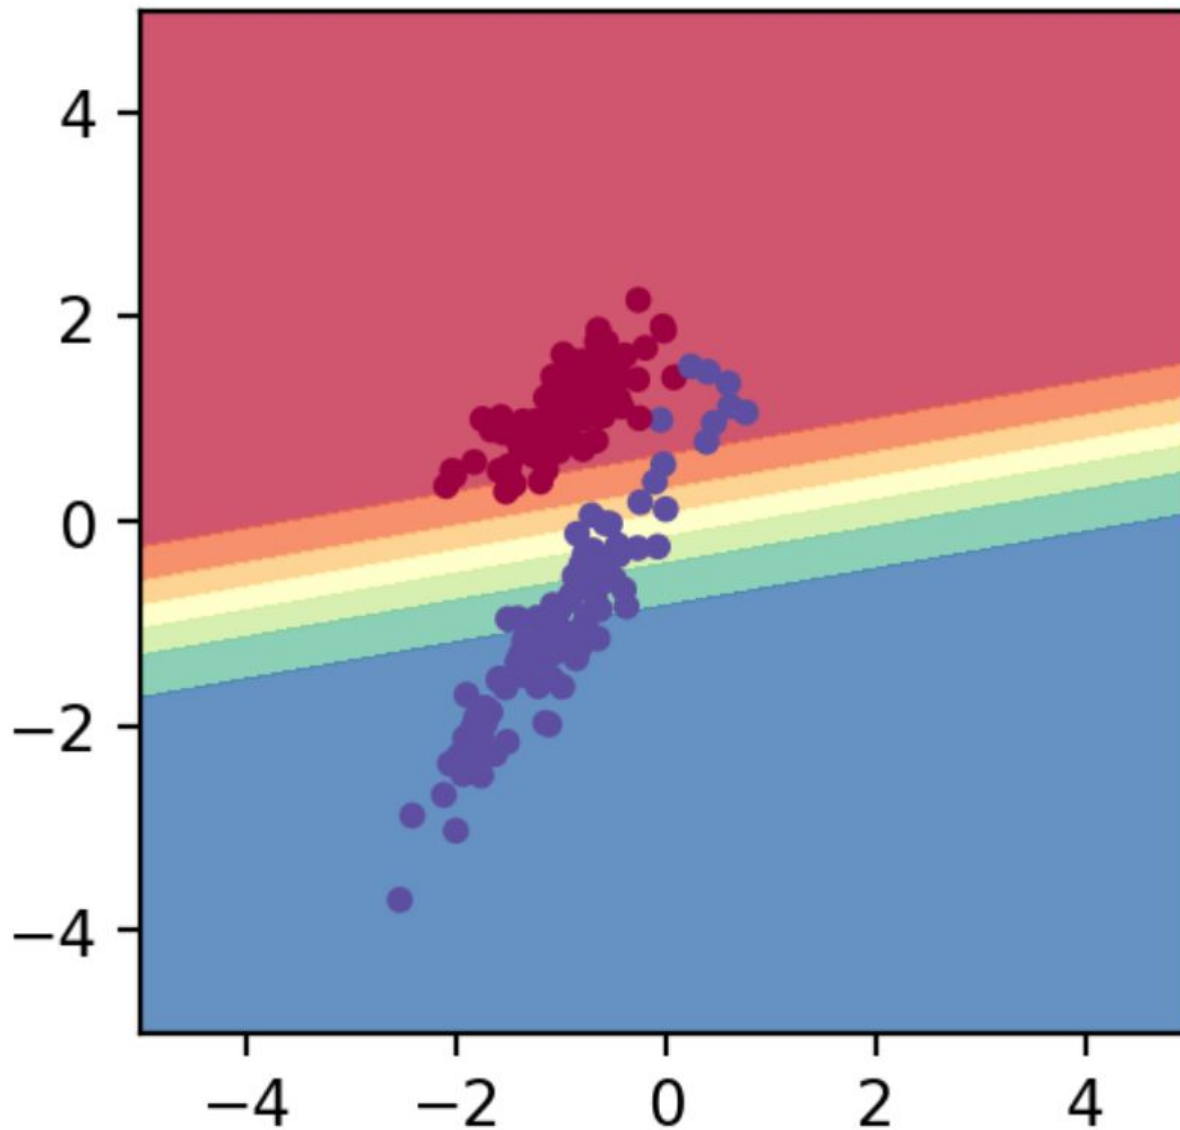
<http://www.inference.org.uk/mackay/itila/book.html>

These slides and classworks follow some of the structure of Chapter 38 & 39.

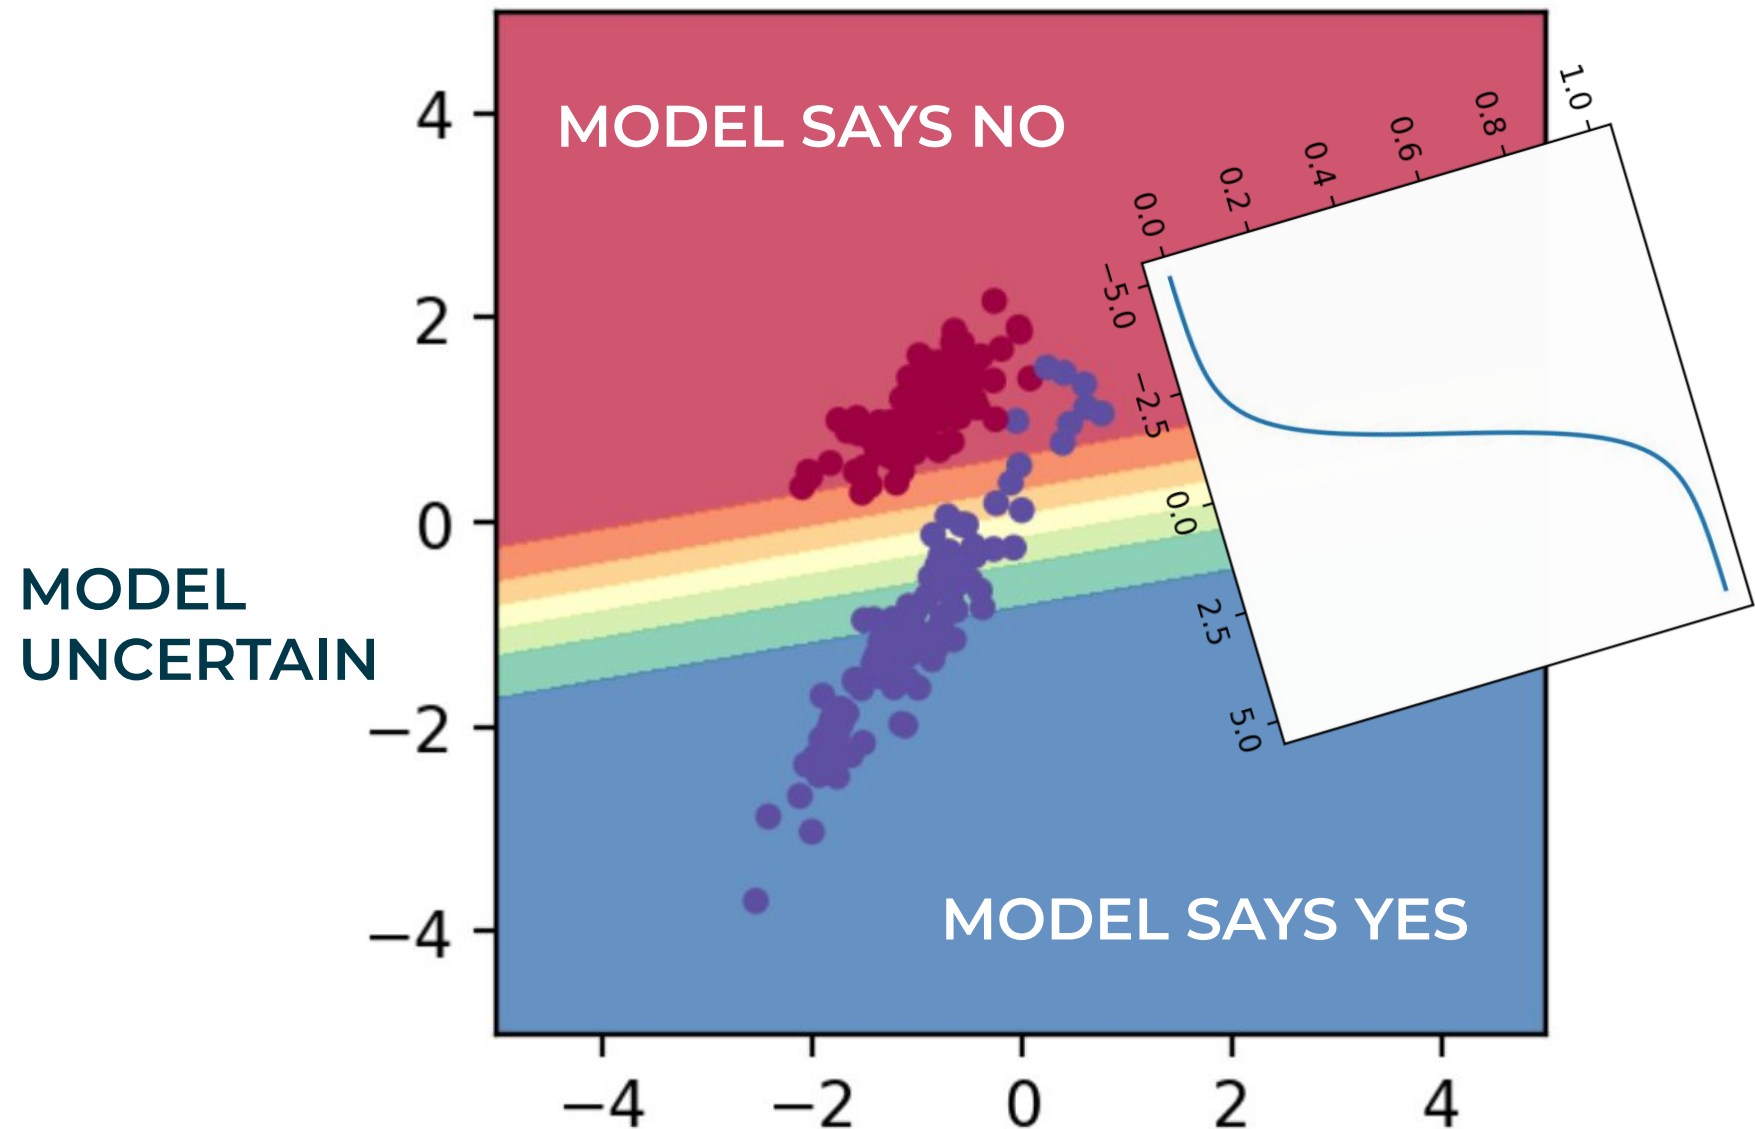


# §3.2 Practical

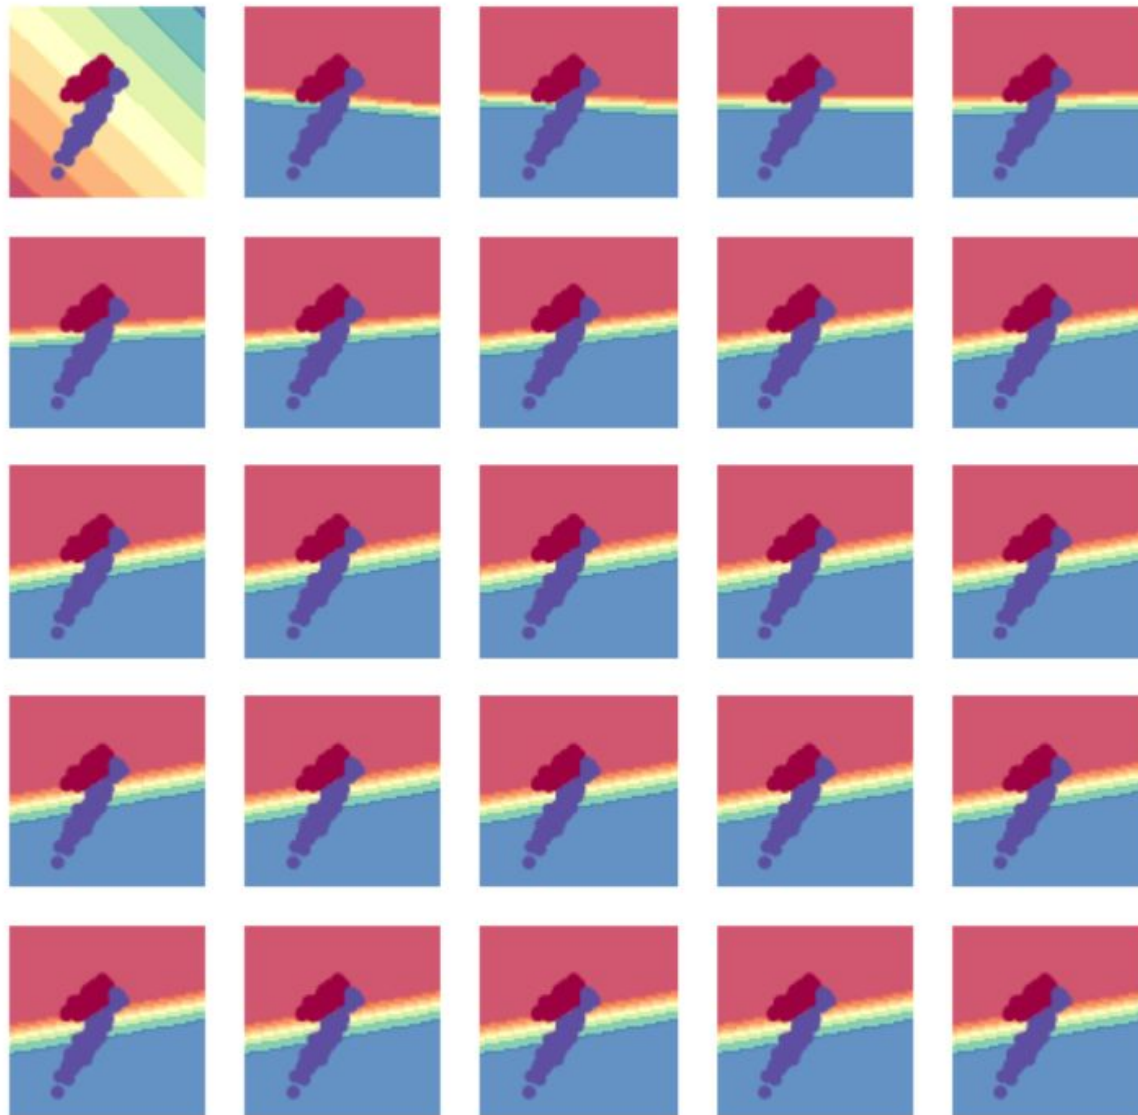
- Add missing code:
  - Sigmoid activation function
  - Neuron function
  - Training loop
- You can now train a model & visualise the decision boundary!
  - Document how the training performs, with different training rates ('eta') and weight decay ('alpha')
  - Compare training with a simple linear classifier
  - How does the neuron fair on the more difficult Gaussian dataset?
- Advanced concepts
  - Change the activation function. (Think about the gradients.)
  - Often neural networks have an additional 'bias' input. Add this to your code.
  - Batch training - currently all data is used to build the single gradient.
  - What happens if you try and use the method for regression (against a function)?
  - Can you improve the regression performance by adding extra neurons side-by-side, each fitting a different part of the function?
  - Compare regression to Gaussian processes:  
[https://jarvist.github.io/2021-PhysicsMachineLearningPracticum/02\\_GaussianProcessPotentialEnergySurface.html](https://jarvist.github.io/2021-PhysicsMachineLearningPracticum/02_GaussianProcessPotentialEnergySurface.html)
- Suggested homework / self-study
  - David MacKay, Information Theory, Inference and Learning Algorithms (ITILA), 2003, Chapters 38–42.
  - PyTorch '60 minute Blitz'  
[https://pytorch.org/tutorials/beginner/deep\\_learning\\_60min\\_blitz.html](https://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html)
  - <https://fluxml.ai/tutorials/2020/09/15/deep-learning-flux.html> - Julia ML library, the tutorial is based on the Pytorch 60 minute Blitz

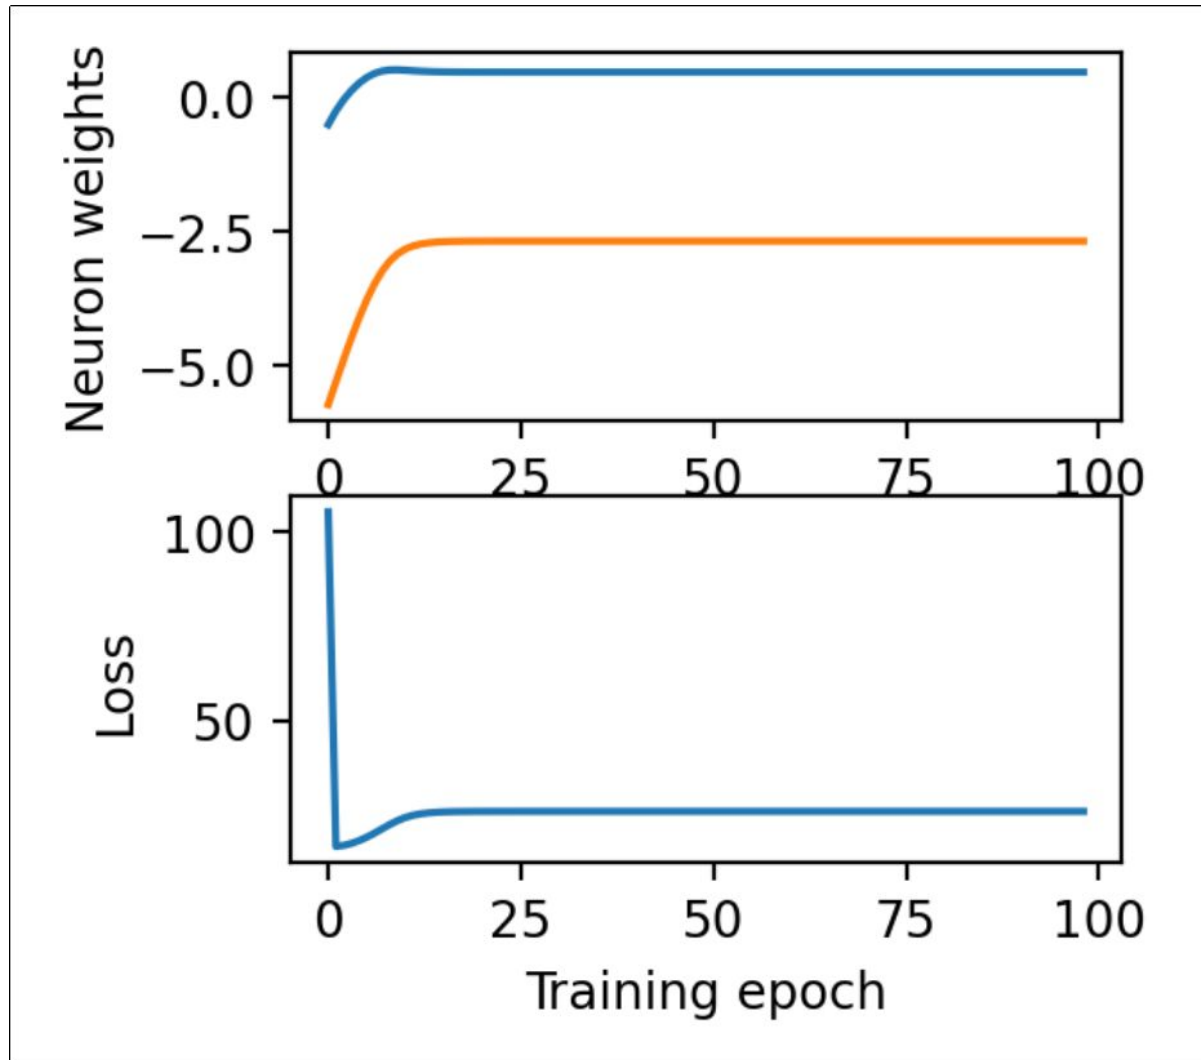


## Decision boundary visualisation



Decision boundary visualisation





## 'Learning curves'



# Observed Values

		Observed Values	
		Positive	Negative
Predicted	Positive	TRUE POSITIVE	FALSE POSITIVE
	Negative	FALSE NEGATIVE	TRUE NEGATIVE

Precision =

$$\frac{TP}{TP + FP}$$

Accuracy =

$$\frac{TP + TN}{\text{Total}}$$

## CONFUSION MATRIX

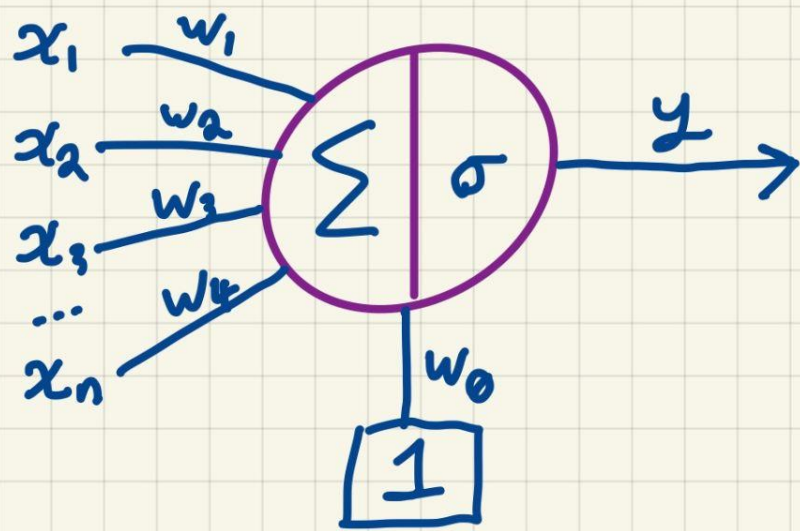
# Observed Values

		Observed Values	
		Positive	Negative
Predicted	Positive	TRUE POSITIVE	FALSE POSITIVE
	Negative	FALSE NEGATIVE	TRUE NEGATIVE

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Accuracy} = \frac{TP + TN}{\text{Total}}$$

CONFUSION  
MATRIX



Training: (supervised)

$$e = y - t$$

$$g_i = -e x_i$$

$$\Delta w_i = -\eta \sum_n g_i$$

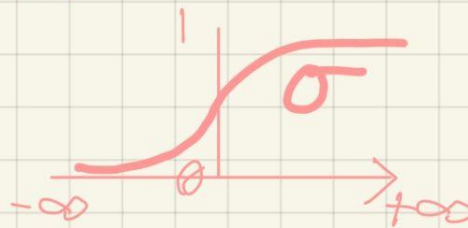
Neuron:

$$a = \sum_i w_i x_i$$

$$y = \sigma(a)$$

Linear

Non-linear

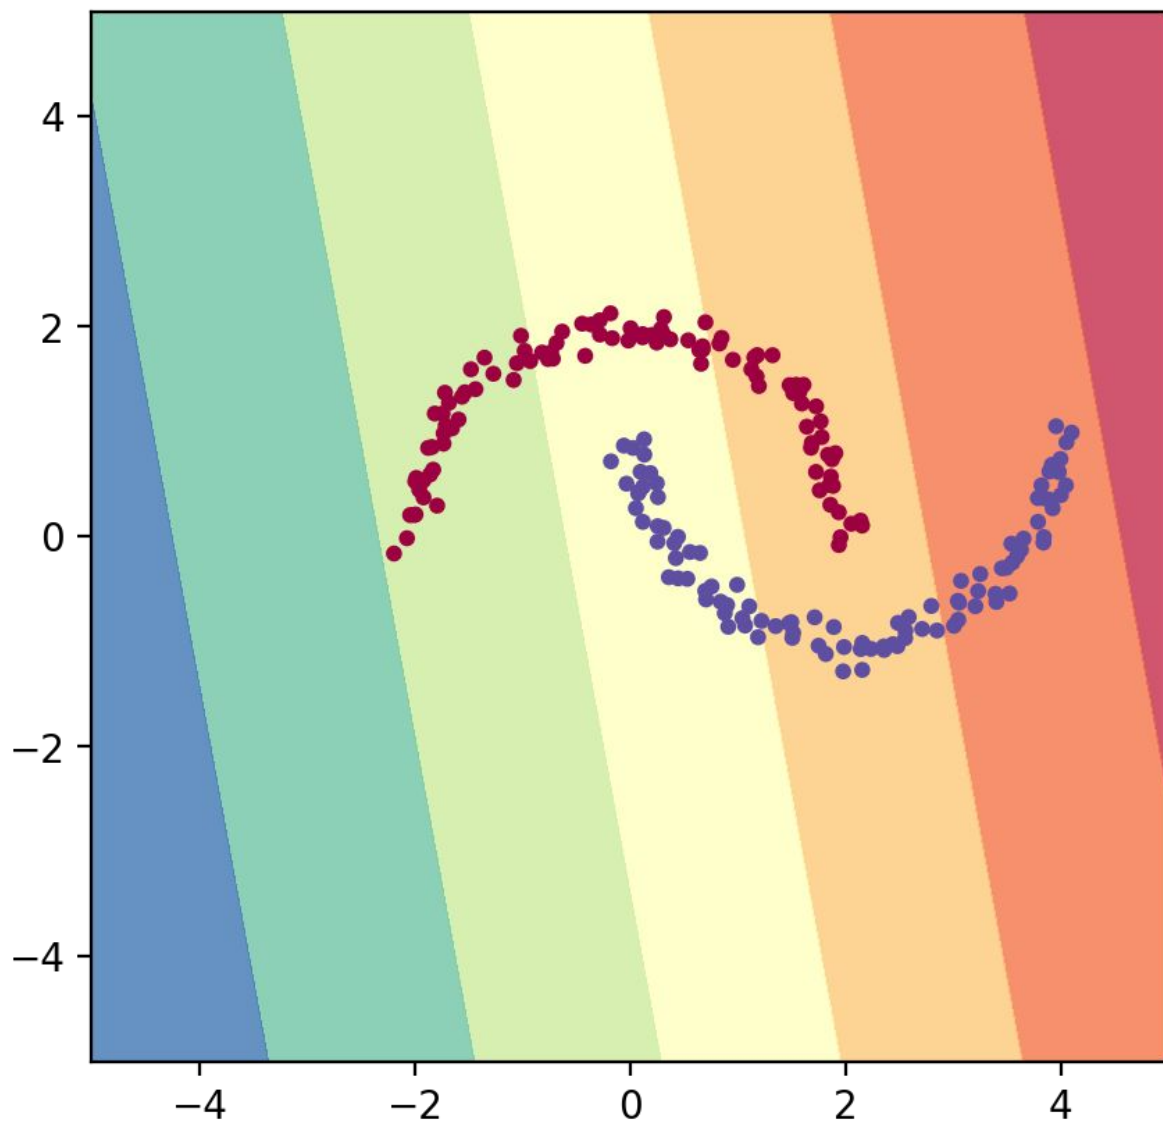


$$a = X * w$$

$x_1$	$x_2$
$x_1$	$x_2$
$x_1$	$x_2$
$x_1$	$x_2$

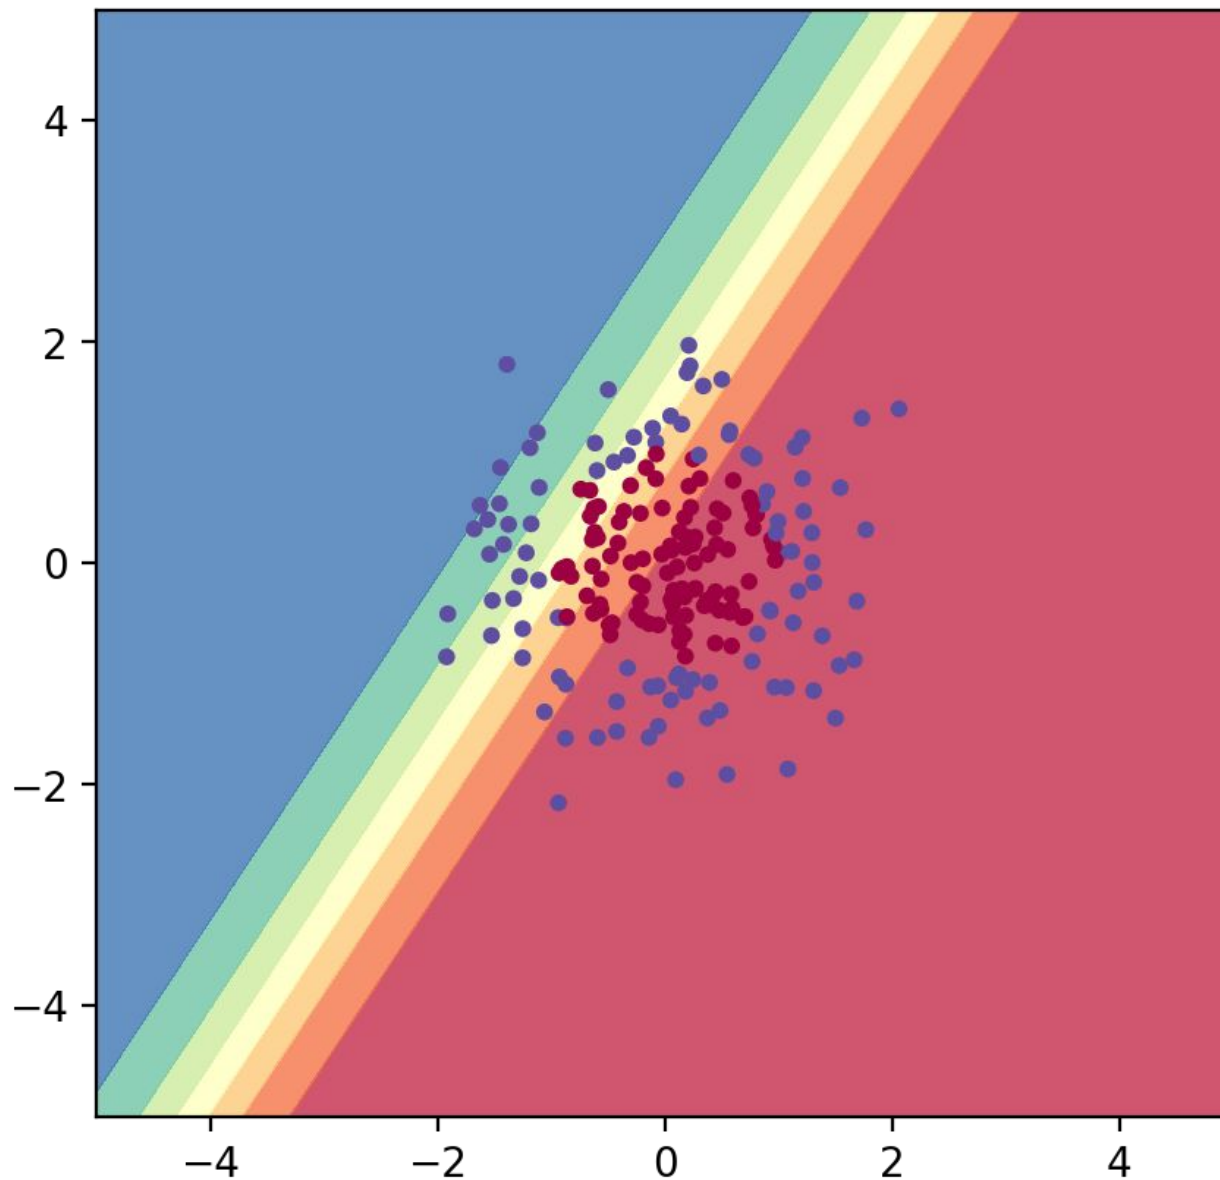
$w_1$	$w_2$
-------	-------

Single neuron by hand. No magic involved!

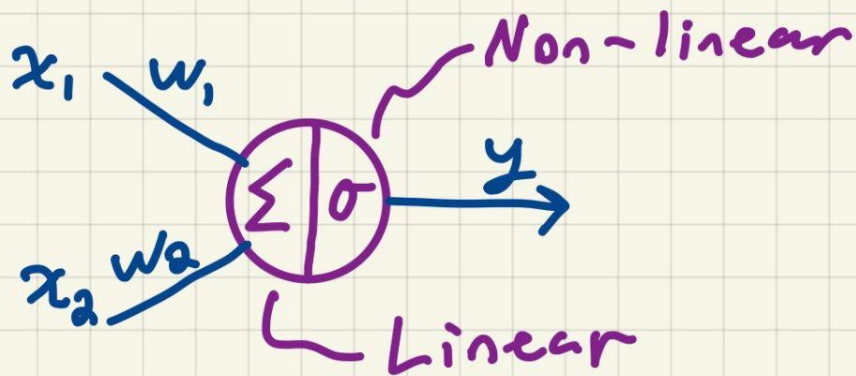




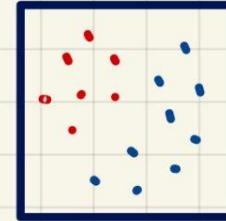




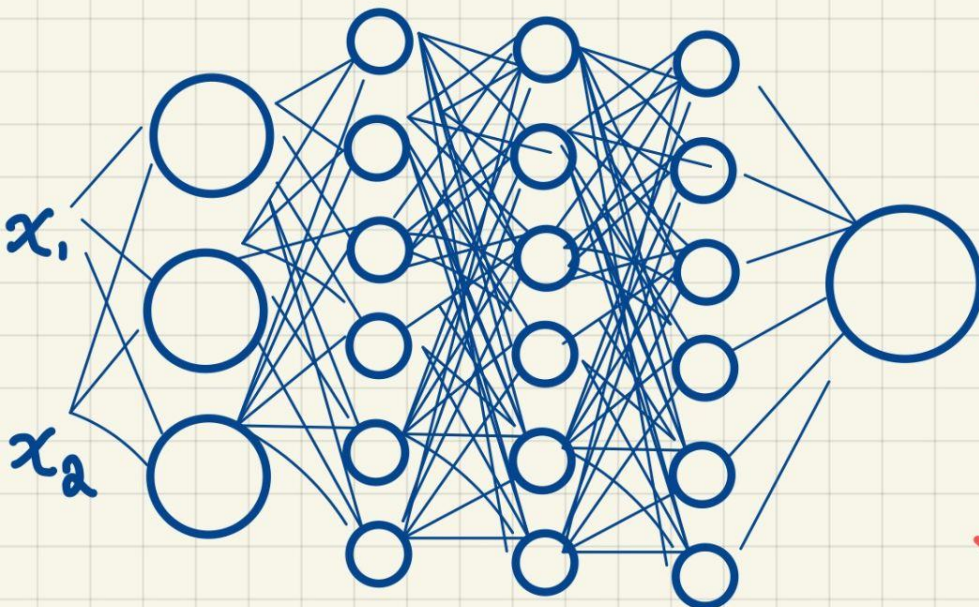
**Similarly poor on overlapping Gaussians...**



- learns from data
- only forms linear decision boundary



Depth

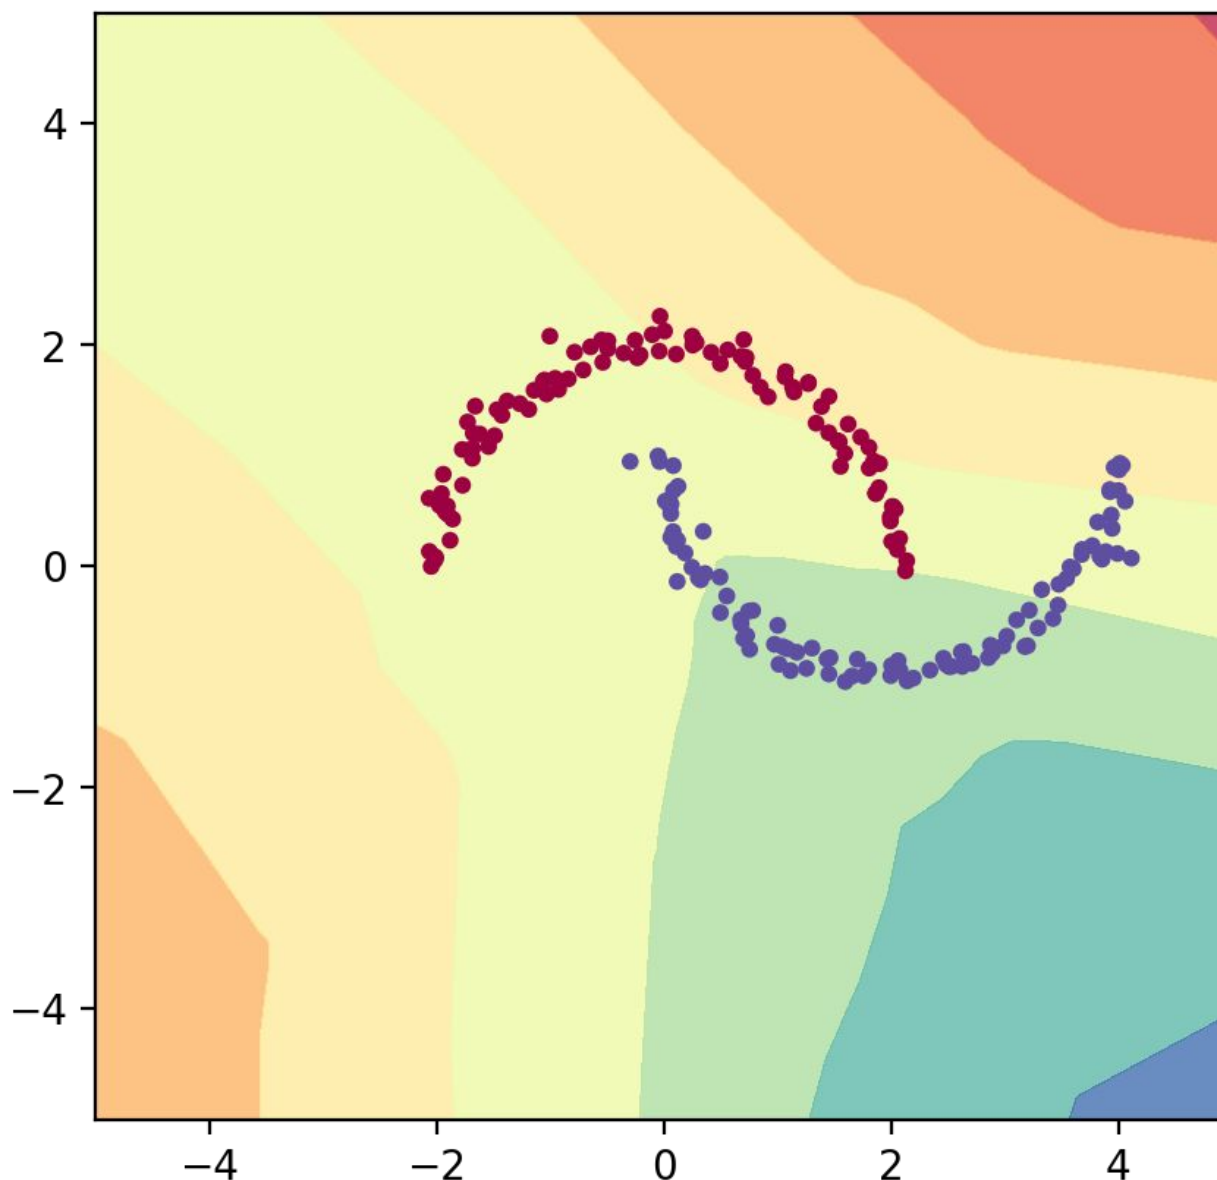


Width

- universal function approximator
- millions of weights trivial
- billions of weights

SOTA

Magic comes from emergent behaviour...

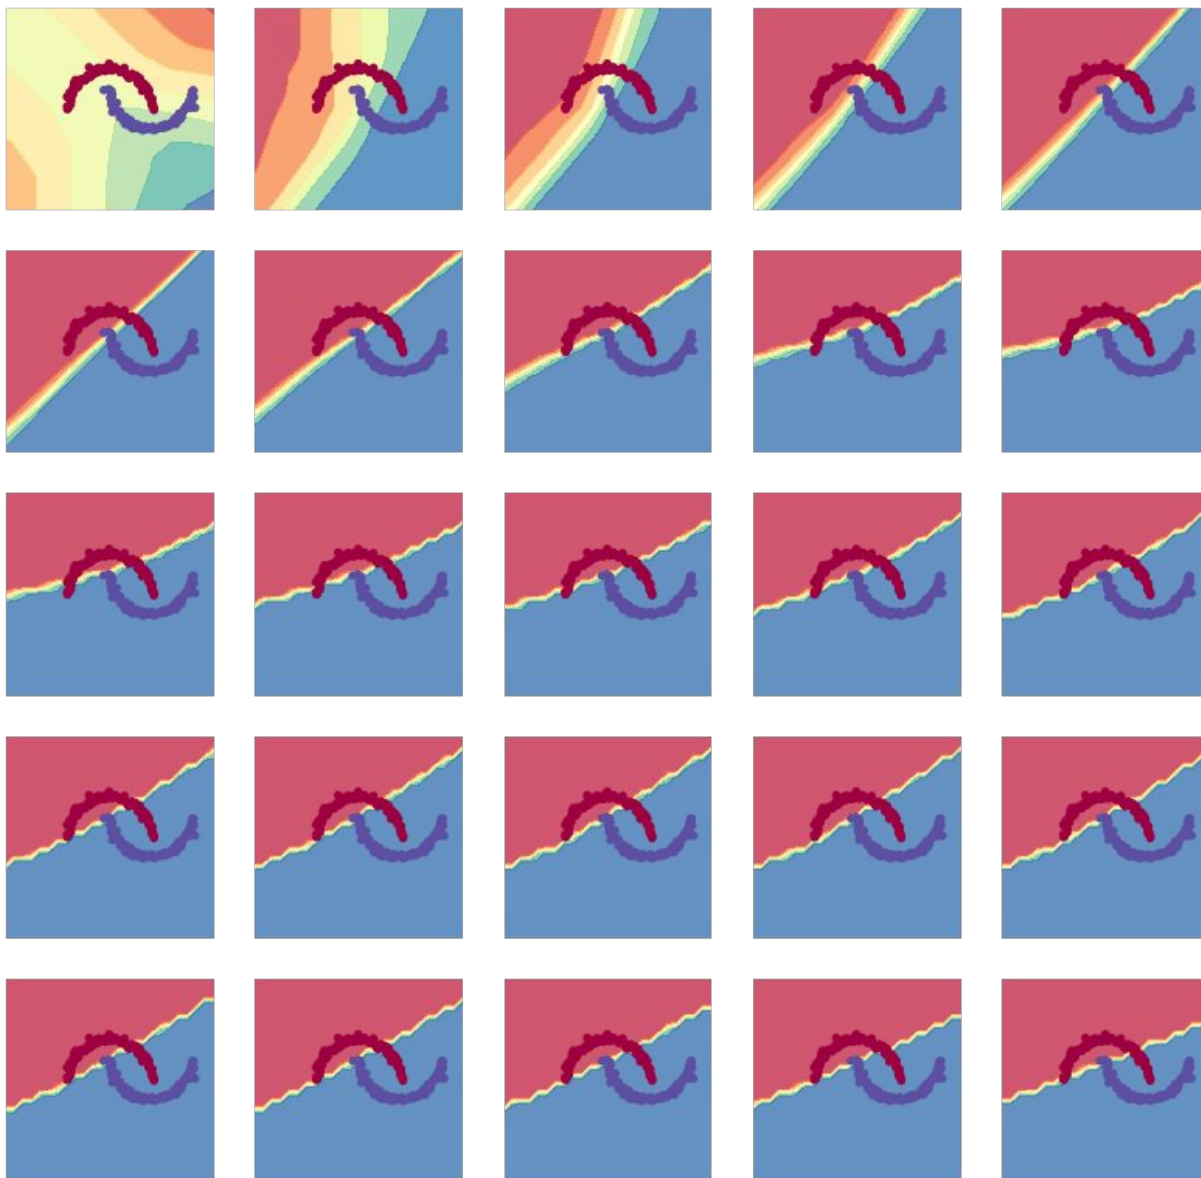


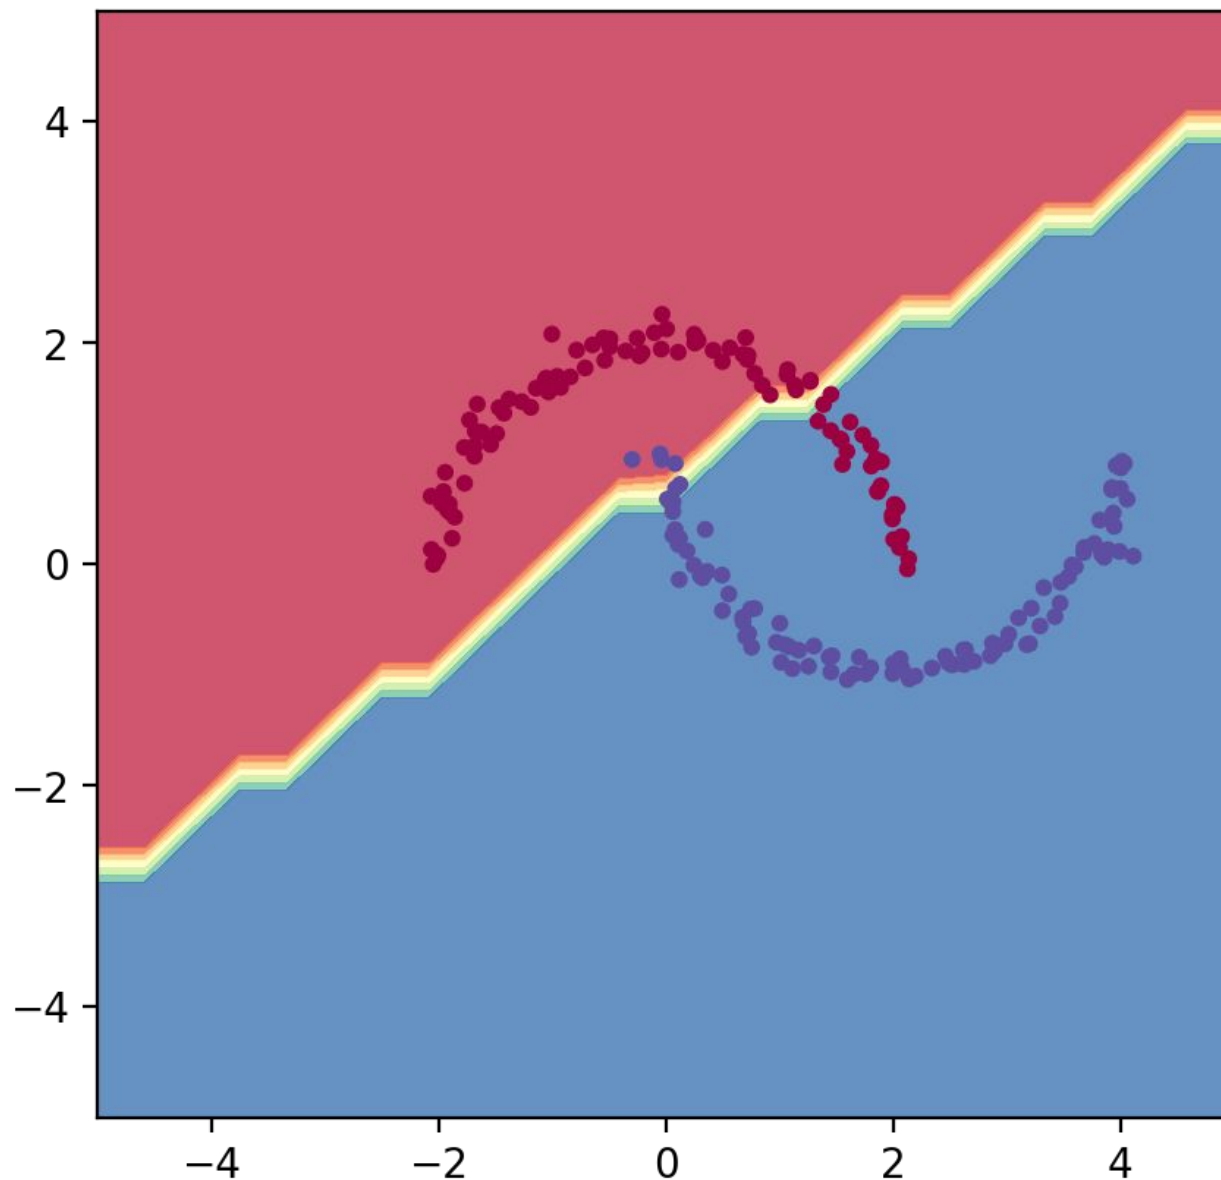
## Multilayer perceptron: 2 $\Rightarrow$ 40, ReLU; 40 $\Rightarrow$ 1, Sigmoid Training by back propagation (G.Hinton 1986)

Rumelhart, D., Hinton, G. & Williams, R. Learning representations by back-propagating errors. Nature 323, 533–536 (1986).

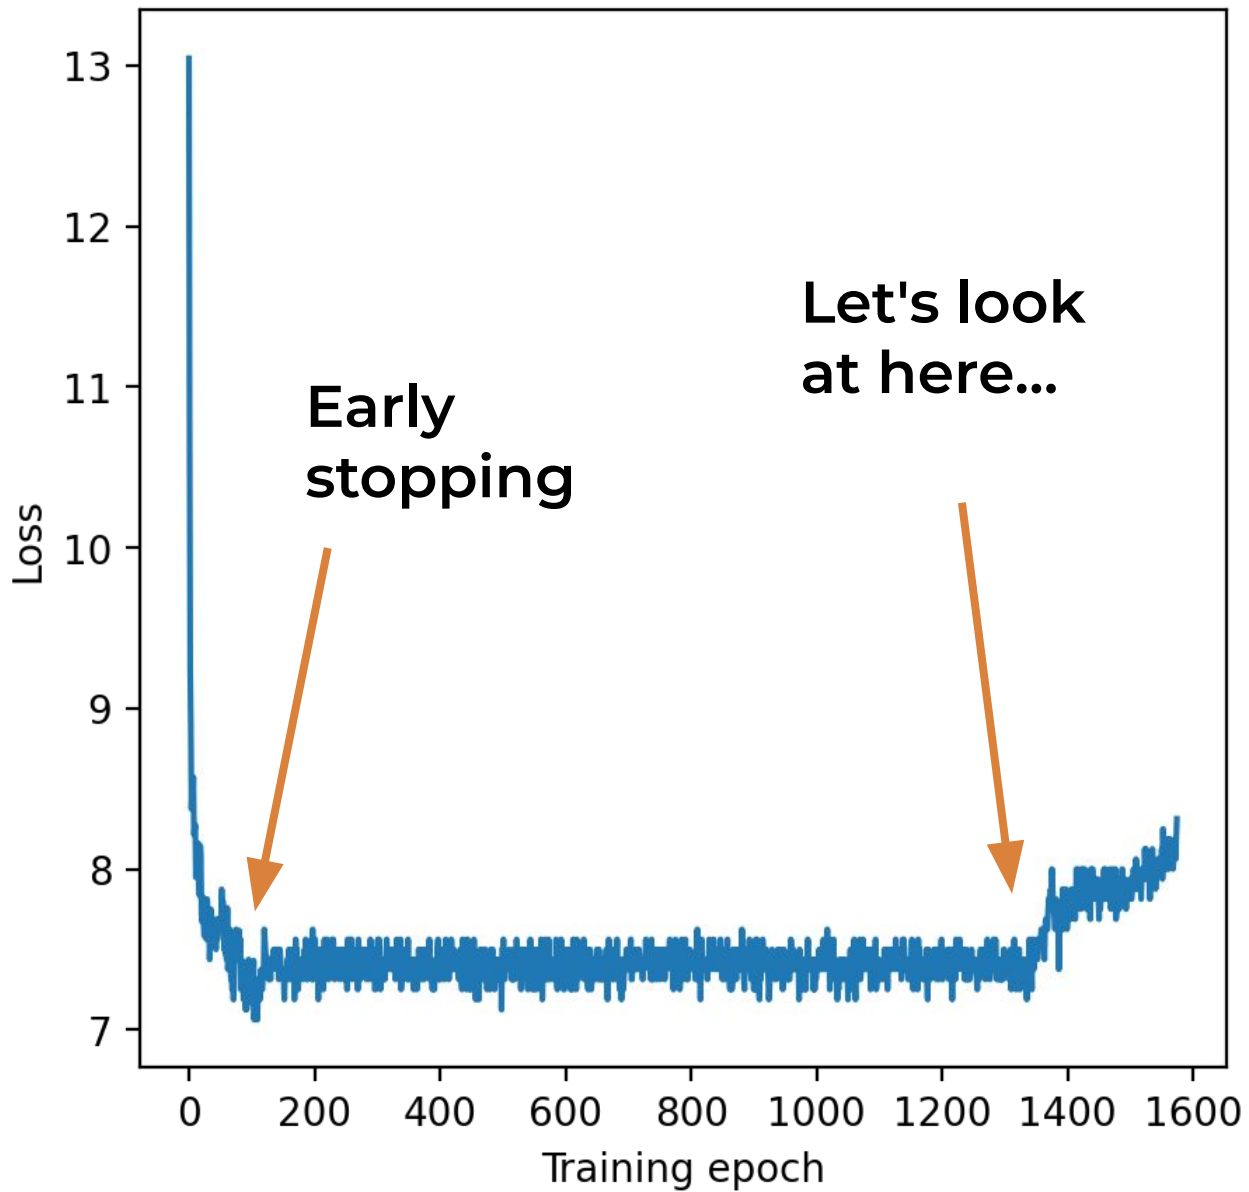
<https://doi.org/10.1038/323533a0>



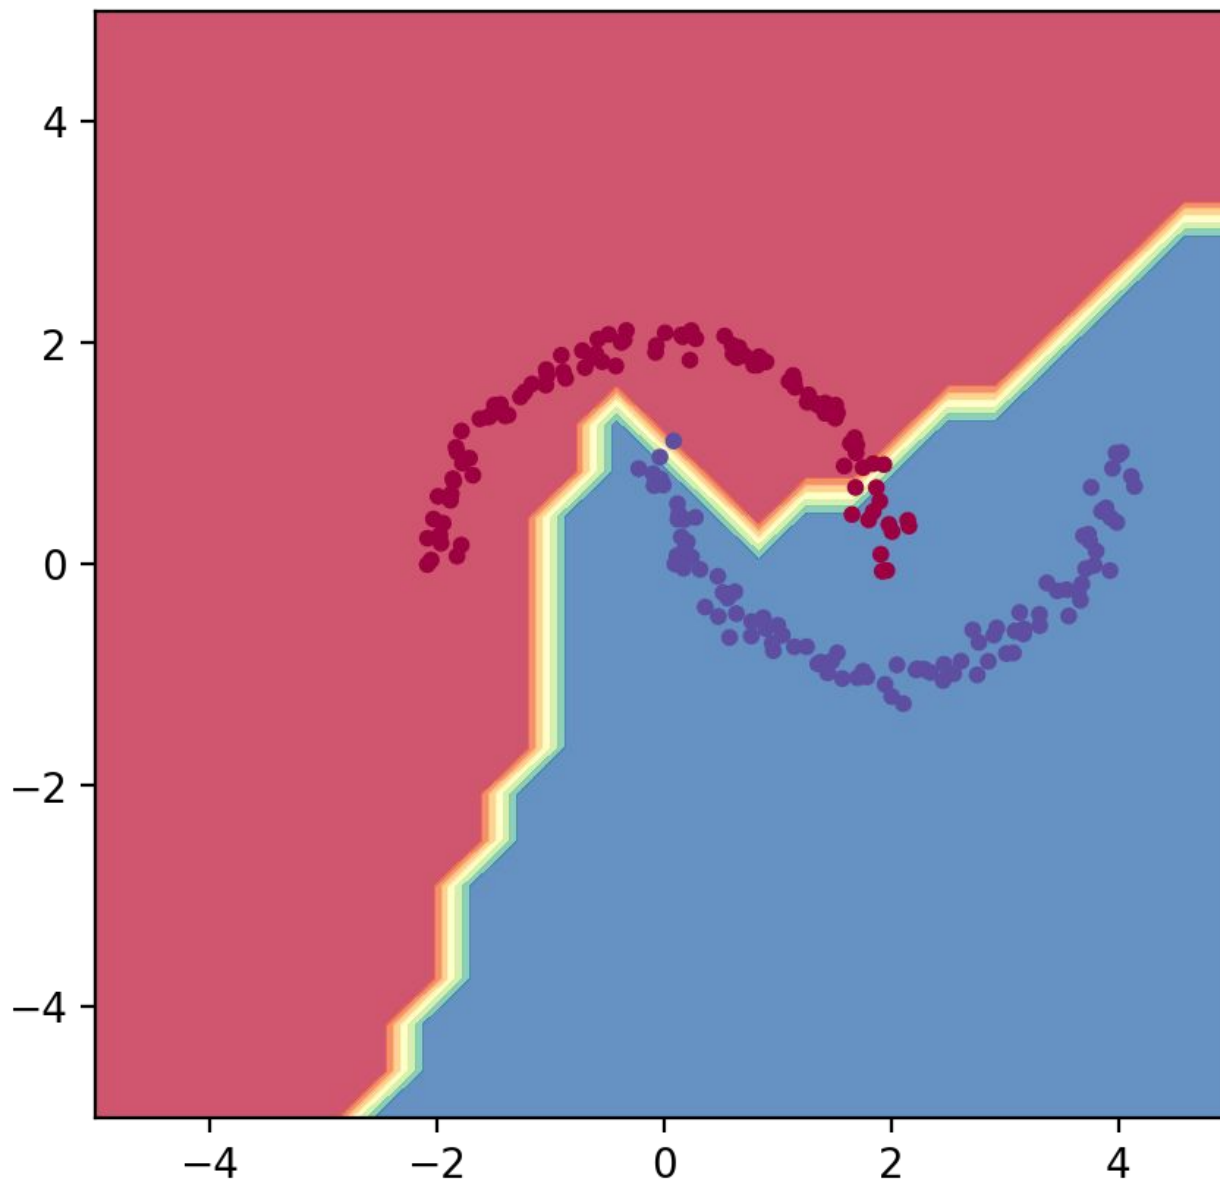




**OK! Solution not so good (accuracy = 35%?)**  
**But non-linear, and interesting wiggles**



**Training is highly empirical; can take a long time.**



# §3.3 Practical

- Add missing (PyTorch!) code:
  - Rebuild a single neuron in Pytorch
    - Layers, loss function
    - accuracy function: true positives + true negatives / total
  - Multilayer Perceptron
    - Define architecture; define loss function
- Single neuron experiments - now with PyTorch
  - Change the activation function. (Think about the gradients.)
  - What effect does the bias have?
  - Set the weights by hand, and document the decision boundary.
- Multi-layer perceptron experiments - with PyTorch
  - Experiment with model construction, task, learning parameters, etc.
    - Document as you go! (Jupyter is really bad for this.)
    - Scoreboard for highest accuracy against each task:
      - Moons
      - Gaussians
      - Your own challenge here?
- Extend PyTorch code to Regression
  - What happens if you try and use the method for regression (against a function)?
    - Collect data points from a function, and fit over a range, i.e.
      - $y=\sin(x)$  ;  $y=x$  ;  $y=x^2$  ;  $y=\sinh(x)$
  - Can you improve the regression performance by adding extra neurons side-by-side, each fitting a different part of the function?
  - Compare regression to Gaussian processes:  
[https://jarvist.github.io/2021-PhysicsMachineLearningPracticum/02\\_GaussianProcessPotentialEnergySurface.html](https://jarvist.github.io/2021-PhysicsMachineLearningPracticum/02_GaussianProcessPotentialEnergySurface.html)